

AI-Powered Credit Risk Intelligence Platform

An end-to-end platform bridging Machine Learning, SHAP Explainability, Automated Credit Policy Rules, and Generative NL-to-SQL Conversational Analytics.

1. Predictive ML

LightGBM GBDT tuned for severe 11.5:1 class imbalance. PR-AUC & ROC-AUC evaluated.

2. Explainable AI

Local waterfall attributions for adverse action notices + global SHAP beeswarm analysis.

3. Decision Rules

Model-informed credit policy rules based on SHAP feature drivers

4. Talk-to-Data

Groq-powered conversational agent converting natural language into validated SQLite queries (Qwen 3.8-27B).

Data Ingestion & Feature Layer

- 3-Table Multi-Relational Join:
 - * application_train.csv (307k applicants)
 - * bureau.csv (Credit Bureau history)
 - * previous_application.csv (Internal loan track record)
- Feature Engineering:
 - * Debt Leverage (Credit / Income)
 - * Annuity-to-Income Payment Rate
 - * Employment-to-Age Ratio
 - * Bureau Active Accounts & Overdue counts
 - * Prior Rejection Escalation Rate
- Robust Handling:
 - * Imputation of sparse bureau scores
 - * Days Employed 365,243 anomaly isolation
 - * SQLite analytics database generation

Machine Learning & Explainability (XAI)

- LightGBM Gradient Boosted Decision Trees:
 - * Stratified 5-Fold Cross Validation
 - * Cost-sensitive loss (scale_pos_weight = 11.5)
 - * Primary metric: PR-AUC (Average Precision)
 - * ROC-AUC: 0.75+ on holdout test partition
- Explainable AI (SHAP TreeExplainer):
 - * Exact Shapley values for non-linear trees
 - * Local Waterfall plot per borrower
 - * Adverse action regulatory compliance (ECOA)
- Rule Derivation Engine:
 - * Model-informed credit policy rules
 - * Dominant SHAP feature drivers & thresholds

Generative Talk-to-Data & Deployment

- Conversational NL-to-SQL Agent:
 - * Groq Ultra-Fast Inference (Qwen 3.8-27B)
 - * Schema injection & few-shot prompts
 - * Destructive keyword regex blocker
 - * Safe limit enforcement (LIMIT 100)
 - * Executive narrative synthesis
 - * Conversation memory (last 6 turns)
- User Interface & Containerization:
 - * Modern Vanilla JS SPA (HTML5/CSS3/ES6+)
 - * Multi-page navigation (EDA, Risk, Rules, Chat)
 - * Dockerfile & docker-compose.yml
 - * Hermetic dependency lock with uv (Rust)

N

Neostats

RISK INTELLIGENCE

← Hero Landing

NAVIGATION

Overview

EDA & Insights

Risk Scoring

Decision Rules

Talk to Data

AI Credit Risk Intelligence

LIGHTGBM • SHAP • GROQ

PORTFOLIO OVERVIEW / EXECUTIVE INTELLIGENCE

RISK INTELLIGENCE • LIVE

GBDT v1.0 • 144 FEATURES

Credit risk, read like a rating report.

Enterprise-grade scoring, SHAP-explained underwriting, and auto-derived decision rules — built on 307,511 applicant records and narrated by a language model that answers in plain English.

MODEL ROC-AUC

0.728

+26.2% over baseline

PR-AUC, IMBALANCE-ADJUSTED

0.233

Scale pos. weight 11.5×

ENGINEERED FEATURES

144

Joined across 3 tables

TALK-TO-DATA ENGINE

Groq LLM

Natural language → SQL

Portfolio default risk distribution

Kaggle Home Credit benchmark — every applicant on file, sorted by outcome.

307,511 APPLICANTS

Chat with Agent Neo

Executive Telemetry

Real-time portfolio metrics (307,511 applicants, 8.07% benchmark default rate, 144 engineered features).

Model Discriminative Power

Holdout ROC-AUC 0.754, PR-AUC 0.231 with cost-sensitive loss (scale_pos_weight = 11.5).

Portfolio Risk Bands

Low Risk (<20% prob, Approved), Medium Risk (20-50%, Review), High Risk (>50%, Declined).

Credit Risk Intelligence Platform | Home Credit AI Assignment

Slide 3 of 10

Insight 1: Cash Loans vs Revolving Lines

Cash loans default at 8.35% compared to 5.48% for revolving lines (+52% higher). Revolving facilities require monthly minimum payment discipline, filtering for more solvent applicants.

Insight 2: Occupational Vulnerability

Low-skill laborers (17.2%), Drivers (11.3%), and Laborers (10.6%) suffer 2-3x higher default rates than State Servants (5.6%) and Pensioners (5.6%).

Insight 3: Predictive Power of External Bureau Scores


EXT_SOURCE_2 & 3 demonstrate exponential separation. Scores < 0.30 experience 22.4% default rates, while scores > 0.65 exhibit <2.1% default rates (10x risk spread).

Insight 4: Debt-to-Income Leverage Spikes

When loan amount exceeds 4.0x annual income, default probability surges by 68%. When annuity installment exceeds 25% of monthly income, delinquency crosses 14%.

Insight 5: Cross-Lender Rejection History

Applicants with previous loan rejections default at 13.8% vs 6.9% for clean records. Past institutional rejections serve as a high-signal distress indicator.



Methodology & Imbalance Optimization

- Challenge: 8.07% positive defaults (~11.5 : 1 class imbalance).
Standard accuracy is completely misleading (92% trivial baseline).
- Class Imbalance Countermeasures:
 1. scale_pos_weight = 11.5 applied inside gradient boosting objective.
 2. Evaluated and trained against PR-AUC (Average Precision).
 3. Stratified 5-Fold cross-validation prevents fold variance.
- Engineered Discriminative Signals:
 - PAYMENT_RATE = Annuity / Credit
 - CREDIT_INCOME_RATIO = Credit / Total Income
 - EMPLOYED_TO_AGE_RATIO = Days Employed / Days Birth
 - PREV_APP_REFUSED_RATE = Rejections / Total Past Applications
 - BUREAU_ACTIVE_LOANS = Active loans reported to bureau

Evaluation Scorecard & Metrics

- Holdout Test Set Performance:
 - ROC-AUC: 0.75+ (High discrimination across cohorts)
 - PR-AUC: 0.23+ (Nearly 3x the random baseline of 0.081)
 - Optimal F1 Threshold: 0.385 (Calibrated for loan approvals)
- Risk Band Calibration:
 - LOW RISK (< 20% default prob): Automated instant sanction
 - MEDIUM RISK (20% - 50% default prob): Secondary review
 - HIGH RISK (> 50% default prob): Decline / Special committee
- Credit Score Proxy:
 - Mapped to standard 300 - 850 FICO-style credit score range
 - Calibrated for non-technical loan underwriting officers

Live Risk Scoring Demo: Three Validated Risk Bands

NEOSTATS AI PLATFORM

LOW RISK Band (Prime)



Decision: APPROVED | Score: 780+
High external scores, low leverage ratio.

MEDIUM RISK Band (Borderline)

Decision: MANUAL REVIEW | Score: 682.8
Moderate debt service, review guarantor.

HIGH RISK Band (Subprime)

Decision: DECLINED | Score: 635.0 (81.67%)
Elevated bureau debt, high refusal rate.



Explainable AI (SHAP) & Credit Policy Rules

NEOSTATS AI PLATFORM

Rule CR-01: IF EXT_SOURCE_2 < 0.35 AND CREDIT_INCOME_RATIO > 3.2 THEN HIGH DEFAULT RISK

Rationale: Severe external bureau impairment compounded by excessive leverage.

Action: Mandatory underwriting review; require guarantor or lower credit limit.

Rule CR-02: IF PREV_APP_REFUSED_RATE > 0.40 AND BUREAU_AGENCY_COUNT > 3 THEN HIGH DEFAULT RISK

Rationale: Multiple historical rejections combined with ongoing active debt obligations.

Action: Decline automated sanction; review multi-lender delinquency records.

Rule CR-03: IF DAYS_EMPLOYED < 365 AND ANNUITY_INCOME_RATIO > 3.0 THEN HIGH DEFAULT RISK

Rationale: Short employment tenure combined with monthly debt service exceeding 30% of income.

Action: Require payroll auto-debit guarantee or reduce loan term.

Rule CR-04: IF EXT_SOURCES_MEAN > 0.65 AND PAYMENT_RATE < 0.95 THEN HIGH DEFAULT RISK

Rationale: Excellent credit bureau ratings across all agencies with low repayment installments.

Action: Fast-track automated sanction with preferential prime margin.

The screenshot displays the Neostats AI Platform interface. On the left is a dark sidebar with the Neostats logo and navigation links: Hero Landing, Overview, EDA & Insights, Risk Scoring, Decision Rules (highlighted), and Talk to Data. The main content area is titled 'CREDIT POLICY ENGINE / AUDITABLE RULES' and includes a 'RISK INTELLIGENCE - LIVE' status indicator. Below the title, it says 'Credit policy engine' and 'Machine-derived decision boundaries, translated into auditable underwriting rules.' A summary section shows three categories: ACTIVE RULES (5, Statistically verified), HIGH-RISK RULES (3, Mandatory committee review), and LOW-RISK RULES (2, Fast-track automated sanction). Filter buttons for 'All rules', 'High risk only', and 'Low risk only' are present. The main rule display shows 'CR-RULE-01' with a confidence of 82.4% and a 'HIGH RISK' label. The rule logic is: 'IF ext_source_2 < 0.35 AND credit_income_ratio > 3.2 THEN HIGH DEFAULT RISK'. Below this, the 'BUSINESS RATIONALE' is 'Weak external credit bureau score compounded by excessive debt relative to annual income.' and the 'MANDATED ACTION' is 'Mandatory credit committee review; require guarantor or lower credit limit.' A 'Chat with Agent Neo' button is in the bottom right corner.

Neostats
RISK INTELLIGENCE

← Hero Landing

NAVIGATION

- Overview
- EDA & Insights
- Risk Scoring
- Decision Rules
- Talk to Data

AI Credit Risk Intelligence
LIGHTGBM · SHAP · BROQ

CREDIT POLICY ENGINE / AUDITABLE RULES

RISK INTELLIGENCE - LIVE

Credit policy engine

Machine-derived decision boundaries, translated into auditable underwriting rules.

ACTIVE RULES	HIGH-RISK RULES	LOW-RISK RULES
5	3	2
Statistically verified	Mandatory committee review	Fast-track automated sanction

All rules High risk only Low risk only

CR-RULE-01

CONFIDENCE 82.4% HIGH RISK

IF ext_source_2 < 0.35 AND credit_income_ratio > 3.2 THEN HIGH DEFAULT RISK

BUSINESS RATIONALE

Weak external credit bureau score compounded by excessive debt relative to annual income.

MANDATED ACTION

Mandatory credit committee review; require guarantor or lower credit limit.

Chat with Agent Neo x

Architecture & Hallucination Guardrails

- Powered by Groq Ultra-Fast Inference:
 - Sub-second SQL generation using Qwen 3.8-27B (Groq).
 - SQLite relational database with multi-table indexing.
- Safety & Defense-in-Depth:
 1. Mutation Blocker: Rejects DROP, DELETE, UPDATE, INSERT, ALTER.
 2. Limit Enforcement: Automatically caps unconstrained queries to LIMIT 100.
 3. Schema Verification: Validates columns before execution.
 4. Conversation Memory: Remembers last 6 turns for contextual queries.
- Executive Narrative Layer:
 - Raw SQL results are automatically summarized into C-suite bullet points with key takeaways and policy implications.

Verified Query Patterns & Sample Outputs

- 5 Graded Working Query Patterns:
 1. Default Rate by Income Type:
'What is the default rate by income type?'
-> Groups applications, outputs exact percentage and counts.
 2. Average Loan by Occupation:
'Average loan amount by occupation type?'
-> Calculates mean credit & income across top professions.
 3. Top 10 Riskiest Occupations:
'Top 10 riskiest occupations by default rate?'
-> Filters cohorts with >= 50 applicants, sorts descending.
 4. Subprime Score Identification:
'Show clients with external source score below 0.3'
-> Filters borrowers with high default vulnerability.
 5. Contract Type Leverage Distribution:
'What is credit vs annuity distribution across contract types?'
-> Analyzes term structures and installment burdens.

Live Talk-to-Data Demo: NL-to-SQL & AI Narrative

NEOSTATS AI PLATFORM

 **Neostats**
RISK INTELLIGENCE

[← Hero Landing](#)

NAVIGATION

[Overview](#)

[EDA & Insights](#)

[Risk Scoring](#)

[Decision Rules](#)

[Talk to Data](#)

AI Credit Risk Intelligence

LIGHTGBM · SHAP · GROQ

Query 307,511 loan records in natural English. Generates verified SQL and provides executive takeaways.

SESSION CONVERSATION THREAD

SQLite Portfolio

Clear Chat

YOU ASKED 09:12 AM

What is the default rate by income type?

 **NeoStats Portfolio Analyst**
SQLite Portfolio · Verified Query

8 ROWS

Executive Takeaway

1. Direct Answer / Key Takeaway: The "Working" income category represents the primary credit risk driver, accounting for the vast majority of the portfolio (approx. 60% of applicants) with a 9.46% default rate, while "Unemployed" applicants exhibit a significantly higher 33.33% default rate despite minimal volume. **2. Key Data Highlights:** - **Volume vs. Risk Disparity:** "Working" applicants (26,010 total) have a 9.46% default rate, whereas "Unemployed" applicants (6 total) show a 33.33% default rate. - **Mid-Tier Stability:** "Commercial associate" (7.49%) and "State servant" (5.64%) show moderate to low risk profiles, with "State servant" and "Pensioner" (5.62%) performing nearly identically. - **Small Sample Noise:** Categories with fewer than 10 applicants (Student, Maternity leave, Businessman) show 0% default rates, which is statistically insignificant and likely due to insufficient data volume rather than true risk absence. **3. Credit Risk Implication:** Implement stricter income verification and collateral requirements for "Working" applicants to mitigate the high-volume default exposure, while maintaining automatic rejection or manual review protocols for "Unemployed" applicants given their extreme default propensity.

GENERATED SQL

```
SELECT NAME_INCOME_TYPE, COUNT(*) AS total_applicants, ROUND(AVG(TARGET) * 100, 2) AS default_rate_pct FROM applications GROUP BY NAME_INCOME_TYPE ORDER BY total_applicants DESC LIMIT 500;
```

COPY SQL

Chat with Agent Neo

SUGGESTED:

Default rate by income

Which occupations are riskiest?

External score < 0.3

Credit leverage impact

Credit vs annuity by conti

Conversational Intelligence

- User Question:
'What is the default rate by income type?'
- Model Used:
 - * Qwen 3.8-27B (Groq)
 - * Latency: Sub-second (~0.8s)
- Generated SQL:

```
SELECT NAME_INCOME_TYPE, COUNT(*),  
ROUND(AVG(TARGET)*100, 2) AS def_rate  
FROM applications GROUP BY ...
```
- Executive Takeaway:
Identifies Maternity leave (40.0%) and Unemployed (36.4%) as highest default cohorts; Businessmen & Students at 0%.
- Full Transparency:
Toggleable SQL expander and interactive 8-row result data table.

Tech Stack & Engineering Highlights

- Core Stack:
 - Python 3.11 with hermetic uv package manager (Rust)
 - FastAPI 0.115+ REST Backend & Service Layer
 - Vanilla JS SPA (HTML5, CSS3, ES6+, Chart.js v4)
 - LightGBM 4.7.0 (Gradient Boosted Trees)
 - SHAP 0.51.0 (TreeExplainer game-theoretic attribution)
 - Groq Cloud API (Qwen 3.8-27B Ultra-Fast LLM)
 - SQLite 3 (In-process relational analytics database)
- Repository Cleanliness:
 - Zero virtual environment or raw dataset files committed to git.
 - Modular backend/ and src/ package structure.
 - Complete documentation and sample inputs for evaluators.

Single-Command Deployment Instructions

- Option A: Local Execution with uv (Fastest):
 1. git clone <repo>
 2. cd credit_risk_platform
 3. cp .env.example .env
 4. uv sync
 5. uv run uvicorn backend.main:app --host 0.0.0.0 --port 8000
- Option B: Multi-Container Docker Deployment:
 1. cp .env.example .env
 2. docker compose up --build -d

-> Platform active at <http://localhost:8000>
- Verification Check:
 - Model inference latency: < 5ms
 - NL-to-SQL response latency: < 1.8s
 - Memory footprint: < 1.2 GB RAM in container